

Captur

Meeting Intelligence Report

Generated on May 18, 2026 at 07:13 PM

Meeting Summary

A team is developing a customer support AI agent.

They discussed the project's architecture, scope, and goals.

The team assigned tasks for Sprint 1 and set latency targets.

They aim to create a warm and efficient agent.

The team will reconvene in two weeks for a review.

Key Insights

Action Items:

- Reza to add input sanitization to the system prompt
- Dana to build the handoff so the whole conversation history gets passed to the rep's dashboard
- Tom to set up evaluations and red-team the agent
- Priya to write a persona brief
- Reza to deliver the core agent loop and start on the eval framework
- Sarah to set up the project board and schedule the Sprint 1 review

Decisions:

- To use a ReAct-style agent loop for the core architecture
- To narrow the scope of tools for v1 to three tools: search the knowledge base, look up order status, and escalate to a human agent
- To prioritize a seamless escalation experience

- Use Claude Sonnet for the main agent loop
- Invest time in creating a persona brief for the agent

Task Owners:

- Reza: Implementing grounding and adding explicit instruction
- Dana: Building the three tool backends and session memory store
- Tom: Building eval harness and compiling the first 100 annotated examples for the eval dataset
- Priya: Writing persona brief and designing spec for the escalation handoff UI and monitoring dashboard
- Sarah: Setting up the project board and scheduling the Sprint 1 review

Deadlines:

- Tom: Compile the first 100 annotated examples for the eval dataset by end of week
- Sarah: Schedule the Sprint 1 review for two weeks out
- Priya: Finish the persona brief and create the design spec this week
- Sarah: Send out the meeting notes and Sprint 1 board link this afternoon